

Problem Set

2024-08-14

Group Members: Quinlan O’Connell, Neal Makwana, Maru Lasala, Ashley Hattendorf

Wrangling Billboard Top 100

Part A: Make a table of the top 10 most popular songs since 1958, as measured by the *total number of weeks that a song spent on the Billboard Top 100*. Note that these data end in week 22 of 2021, so the most popular songs of 2021 will not have up-to-the-minute data; please send our apologies to The Weeknd.

Your table should have **10 rows** and **3 columns**: **performer**, **song**, and **count**, where **count** represents the number of weeks that song appeared in the Billboard Top 100. Make sure the entries are sorted in descending order of the **count** variable, so that the more popular songs appear at the top of the table. Give your table a short caption describing what is shown in the table.

```
top_10 <- billboard %>%
  select(performer, song) %>%
  group_by(performer, song) %>%
  summarize(count = n(), .groups = 'drop') %>% #use summarize to count entries per song, .groups = 'drop'
  arrange(desc(count)) %>%
  head(10)
top_10
```

```
## # A tibble: 10 x 3
##   performer      song      count
##   <chr>         <chr>    <int>
## 1 Imagine Dragons Radioactive      87
## 2 AWOLNATION     Sail        79
## 3 Jason Mraz     I’m Yours     76
## 4 The Weeknd     Blinding Lights 76
## 5 LeAnn Rimes    How Do I Live   69
## 6 LMFAO Featuring Lauren Bennett & GoonRock Party Rock Anthem 68
## 7 OneRepublic    Counting Stars  68
## 8 Adele          Rolling In The Deep 65
## 9 Jewel          Foolish Games/You Were Meant~ 65
## 10 Carrie Underwood Before He Cheats 64
```

The table above shows the top 10 most popular billboard songs. “Radioactive” has the highest count, with a count of 87. Other popular songs like Sail, I’m Yours, and Blinding Lights follow, with a count in the 76-79 range.

Part B: Is the “musical diversity” of the Billboard Top 100 changing over time? Let’s find out. We’ll measure the musical diversity of given year as *the number of unique songs that appeared in the Billboard Top 100 that year*. Make a line graph that plots this measure of musical diversity over the years. The x axis should show the year, while the y axis should show the number of unique songs appearing at any position on the Billboard Top 100 chart in any week that year. For this part, please filter the data set so that it excludes


the years 1958 and 2021, since we do not have complete data on either of those years. Give the figure an informative caption in which you explain what is shown in the figure and comment on any interesting trends you see.

```
unique_billboard <- billboard %>%  
  filter(year != 1958 & year != 2021) %>%  
  group_by(year, performer, song) %>%  
  summarize(count = n()) %>%  
  group_by(year) %>%  
  summarize(new_count = n())  
unique_billboard
```

```
## # A tibble: 62 x 2  
##   year new_count  
##   <int>   <int>  
## 1 1959     663  
## 2 1960     700  
## 3 1961     779  
## 4 1962     768  
## 5 1963     754  
## 6 1964     811  
## 7 1965     800  
## 8 1966     832  
## 9 1967     827  
## 10 1968     772  
## # i 52 more rows
```

```
ggplot(unique_billboard) +  
  geom_line(aes(x= year, y= new_count)) +  
  labs(x= "Year",  
       y="Number of Unique Songs",  
       title= "Unique Songs in Billboard Top 100 Each Year")
```

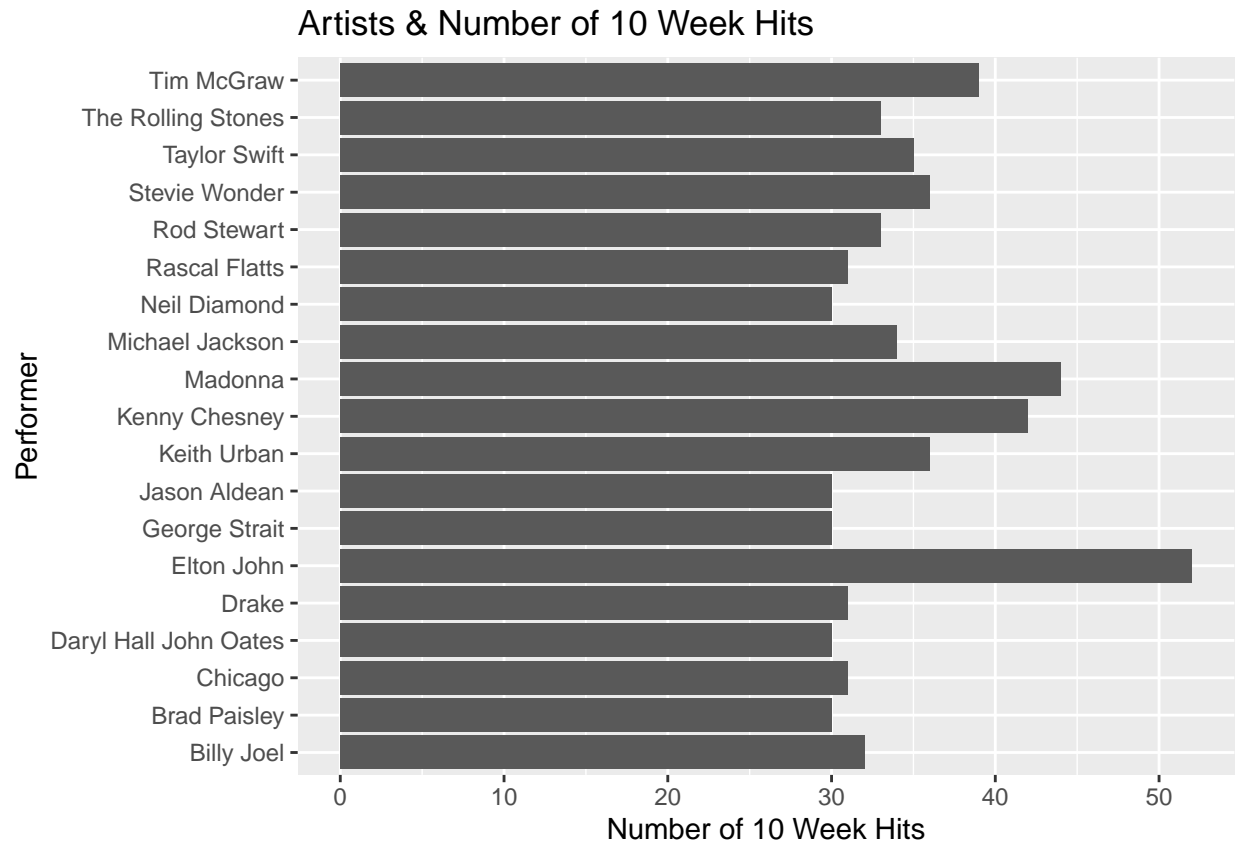
Unique Songs in Billboard Top 100 Each Year



The line graph above shows the number of unique songs that appeared in the Billboard Top 100, based on the year. We can see that the year with the most unique songs to make the list was in about 1967, where the data peaks at about 850 unique songs making the list that year. The data then trends downward until it reached a low in around 2001, at about 380 unique songs on the list. From here, the data trends upward until 2011, where it hits a short dip and then continues back upward until 2020.

Part C: Let's define a "ten-week hit" as a single song that appeared on the Billboard Top 100 for at least ten weeks. There are 19 artists in U.S. musical history since 1958 who have had *at least 30 songs* that were "ten-week hits." Make a bar plot for these 19 artists, showing how many ten-week hits each one had in their musical career. Give the plot an informative caption in which you explain what is shown.

```
ten_weeks_billboard <- billboard %>%
  group_by(performer, song) %>%
  summarize(n = n()) %>%
  filter(n >= 10) %>%
  group_by(performer) %>%
  summarize(n = n()) %>%
  arrange(desc(n)) %>%
  filter(n >= 30)
ggplot(ten_weeks_billboard) +
  geom_col(aes(x = performer, y = n)) +
  labs(x="Performer",
       y="Number of 10 Week Hits",
       title= "Artists & Number of 10 Week Hits") +
  coord_flip()
```




The bar plot above shows artists who have had at least 30 songs that were “ten-week hits” (songs that appeared on the Billboard Top 100 for at least 10 weeks). 19 artists in history have had at least 30 songs that were “ten-week hits”, and we can see above the number of these (x-axis) based on the artist (y-axis). As seen from the data, the artist with the most “ten-week hits” is Elton John, as he has 52 songs that spent at least 10 weeks on the Billboard Top 100. He is followed by Madonna and Kenny Chesney, who have 44 and 42 ten week hits, respectively.

CapMetro Storytelling

Your task is to create a figure, or set of related figures, that tell an interesting story about Capital Metro ridership patterns around the UT-Austin campus during the semester in question. Provide a clear annotation/caption for each figure, but the figure(s) should be more or less stand-alone, in that you shouldn’t need many, many paragraphs to convey its meaning. Rather, the figure together with a concise caption should speak for itself as far as possible.


```
# named it capmetro_UT_data
capmetro_UT = read.csv("capmetro_UT_data.csv")
boarding_data = mutate(capmetro_UT,
                        day_of_week = factor(day_of_week,
                                             levels=c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun")),
                        month = factor(month,
                                       levels=c("Sep", "Oct", "Nov")))
boarding_data = boarding_data %>%
  group_by(hour_of_day, day_of_week, month) %>%
  summarize(avg_board = mean(boarding))
```

```
ggplot(boarding_data) +
  geom_line(aes(x=hour_of_day, y=avg_board, color = month)) +
  facet_wrap(~day_of_week) +
  labs(x="Hour of Day",
       y="Average Boarding",
       title="Average Boarding by Hour of the Day Each Month")
```

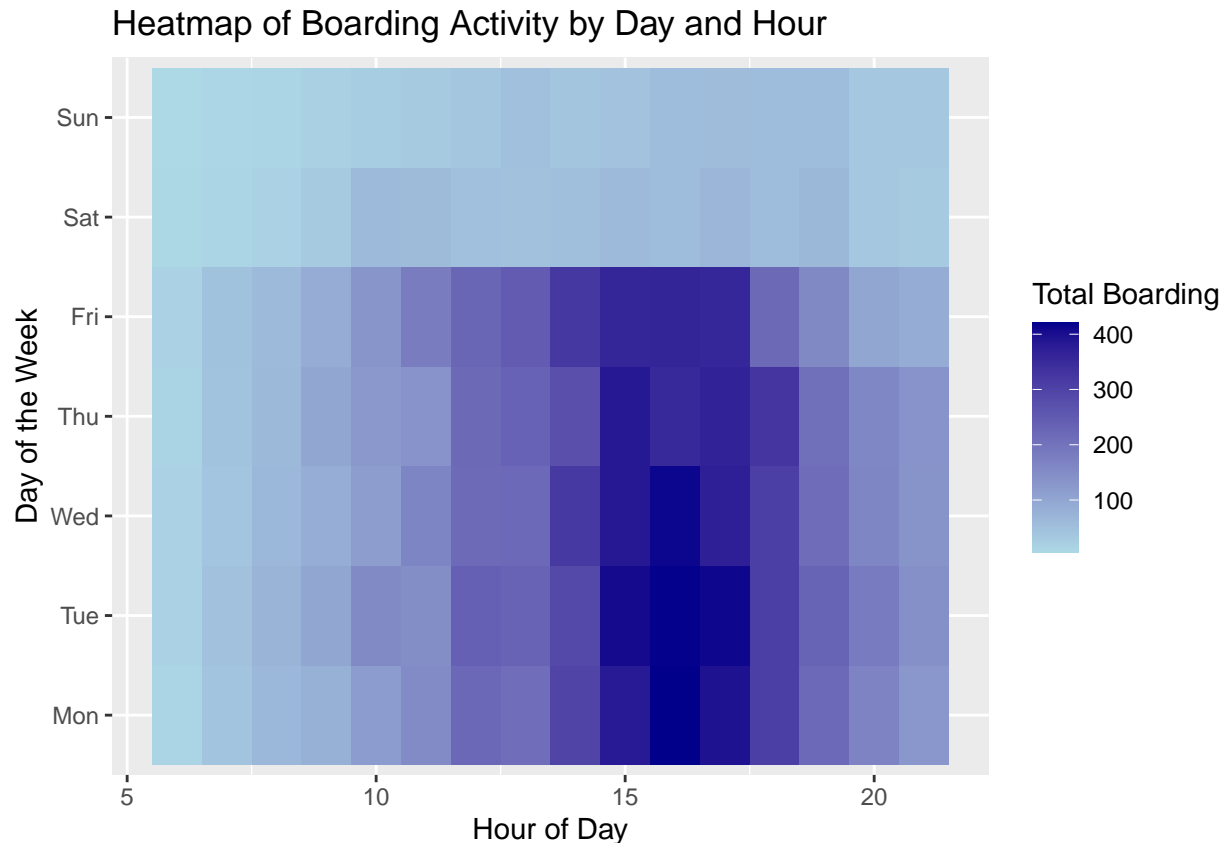


```
temperature_capmetro = capmetro_UT %>%
  group_by(temperature, boarding, hour_of_day, weekend)
ggplot(temperature_capmetro) +
  geom_point(aes(x=temperature, y=boarding, color=weekend), size=0.2) +
  facet_wrap(~hour_of_day) +
  scale_color_brewer(type="qual") +
  labs(x="Temperature",
       y="Boarding",
       title= "Boarding by Temperature on Weekends vs Weekdays")
```

Boarding by Temperature on Weekends vs Weekdays



```
boarding_heatmap <- boarding_data %>%
  group_by(day_of_week, hour_of_day) %>%
  summarize(total_boarding = sum(avg_board))
ggplot(boarding_heatmap, aes(x=hour_of_day, y=day_of_week, fill=total_boarding)) +
  geom_tile() +
  scale_fill_gradient(low="lightblue", high="darkblue") +
  labs(x="Hour of Day",
       y="Day of the Week",
       fill="Total Boarding",
       title="Heatmap of Boarding Activity by Day and Hour")
```



In the top figure, the average boardings on the UT Metro, organized by the day of the week, hour of the day, and month of the year. Three different lines on each graph represent different months of the year, and the X-Axis shows hour of the day while the Y-Axis shows the number of average boardings. According to the data above, the hour of peak boarding is very similar across all three months, but only on the days Monday-Friday. The number of average boardings tends to peak around hour 15 (3PM) during the week, which is likely explained by this time being a heavy travel time for students going to and from class.

In the middle figure, the scatter plot shows the number of UT Metro boardings throughout the hours of the day, based on the temperature and type of day it is (weekday or weekend). It seems like the peak demand for the UT metro is in hours 15-17 (3-5PM) during the week, and 17-18 (5-6PM) on the weekend. Boarding is much less abundant in the early hours (hour 6-9), which is unsurprising considering many college students would be asleep then. According to the data, it seems like when we hold hour of day and weekend status constant, temperature does not seem to have a noticeable effect on the number of UT students riding the bus. When analyzing the data, there is always a heavy distribution of boardings across many different temperatures.

In the bottom figure, the heatmap helps to visualize activity by the hour of the day. We can see a hotter portion on the heatmap in the later afternoon day, which likely indicates the end of the school/workday. We can also see a clear distinction between the weekdays and weekend, which may indicate a decline in ridership when people are not working/in school/commuting to campus.

Rule Mining Association

```
# Generating association rules
rules <- apriori(transactions, parameter = list(supp = 0.001, conf = 0.5))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.5      0.1      1 none FALSE          TRUE      5  0.001      1
## maxlen target  ext
##      10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 9
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [157 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 done [0.02s].
## writing ... [5668 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
# chose base items
```

```
# number of items in rules
length(rules)
```

```
## [1] 5668
```

```
# HIGH Confidence Rules!
#inspect(sort(rules, by = "confidence"))
#arules::inspect(sort(rules, by = "confidence"))

# transactions with whole milk / other vegetables seem to have very high confidence values

# Inspect the top rules by lift
# Sort the rules by lift to find the most interesting ones
sorted_rules <- sort(rules, by = "lift")
# show rules
arules::inspect(sorted_rules[1:30])
```

	lhs	rhs	support	confidence	coverage	lift
## [1]	{Instant food products,	=> {hamburger meat}	0.001220132	0.6315789	0.001931876	18.995654
##	soda}					
## [2]	{popcorn,	=> {salty snack}	0.001220132	0.6315789	0.001931876	16.697793
##	soda}					
## [3]	{baking powder,	=> {sugar}	0.001016777	0.5555556	0.001830198	16.408075
##	flour}					
## [4]	{ham,	=> {white bread}	0.001931876	0.6333333	0.003050330	15.045491
##	processed cheese}					
## [5]	{Instant food products,	=> {hamburger meat}	0.001525165	0.5000000	0.003050330	15.038226
##	whole milk}					

## [6]	{curd, other vegetables, whipped/sour cream, yogurt}	=> {cream cheese}	0.001016777	0.5882353	0.001728521	14.834087
## [7]	{domestic eggs, processed cheese}	=> {white bread}	0.001118454	0.5238095	0.002135231	12.443639
## [8]	{other vegetables, tropical fruit, white bread, yogurt}	=> {butter}	0.001016777	0.6666667	0.001525165	12.030581
## [9]	{hamburger meat, whipped/sour cream, yogurt}	=> {butter}	0.001016777	0.6250000	0.001626843	11.278670
## [10]	{domestic eggs, other vegetables, tropical fruit, whole milk, yogurt}	=> {butter}	0.001016777	0.6250000	0.001626843	11.278670
## [11]	{liquor, red/blush wine}	=> {bottled beer}	0.001931876	0.9047619	0.002135231	11.235269
## [12]	{cream cheese, other vegetables, whipped/sour cream, yogurt}	=> {curd}	0.001016777	0.5882353	0.001728521	11.040638
## [13]	{hard cheese, whipped/sour cream, yogurt}	=> {butter}	0.001016777	0.5882353	0.001728521	10.615219
## [14]	{other vegetables, rolls/buns, root vegetables, tropical fruit, whole milk}	=> {beef}	0.001118454	0.5500000	0.002033554	10.483043
## [15]	{sliced cheese, tropical fruit, whole milk, yogurt}	=> {butter}	0.001016777	0.5555556	0.001830198	10.025484
## [16]	{butter, other vegetables, sugar}	=> {whipped/sour cream}	0.001016777	0.7142857	0.001423488	9.964539
## [17]	{hard cheese, whipped/sour cream, whole milk}	=> {butter}	0.001423488	0.5384615	0.002643620	9.717008
## [18]	{domestic eggs, hard cheese, other vegetables}	=> {butter}	0.001016777	0.5263158	0.001931876	9.497827
## [19]	{fruit/vegetable juice, other vegetables, tropical fruit, whipped/sour cream}	=> {butter}	0.001016777	0.5263158	0.001931876	9.497827
## [20]	{onions, tropical fruit, yogurt}	=> {butter}	0.001118454	0.5238095	0.002135231	9.452599
## [21]	{domestic eggs, other vegetables,					

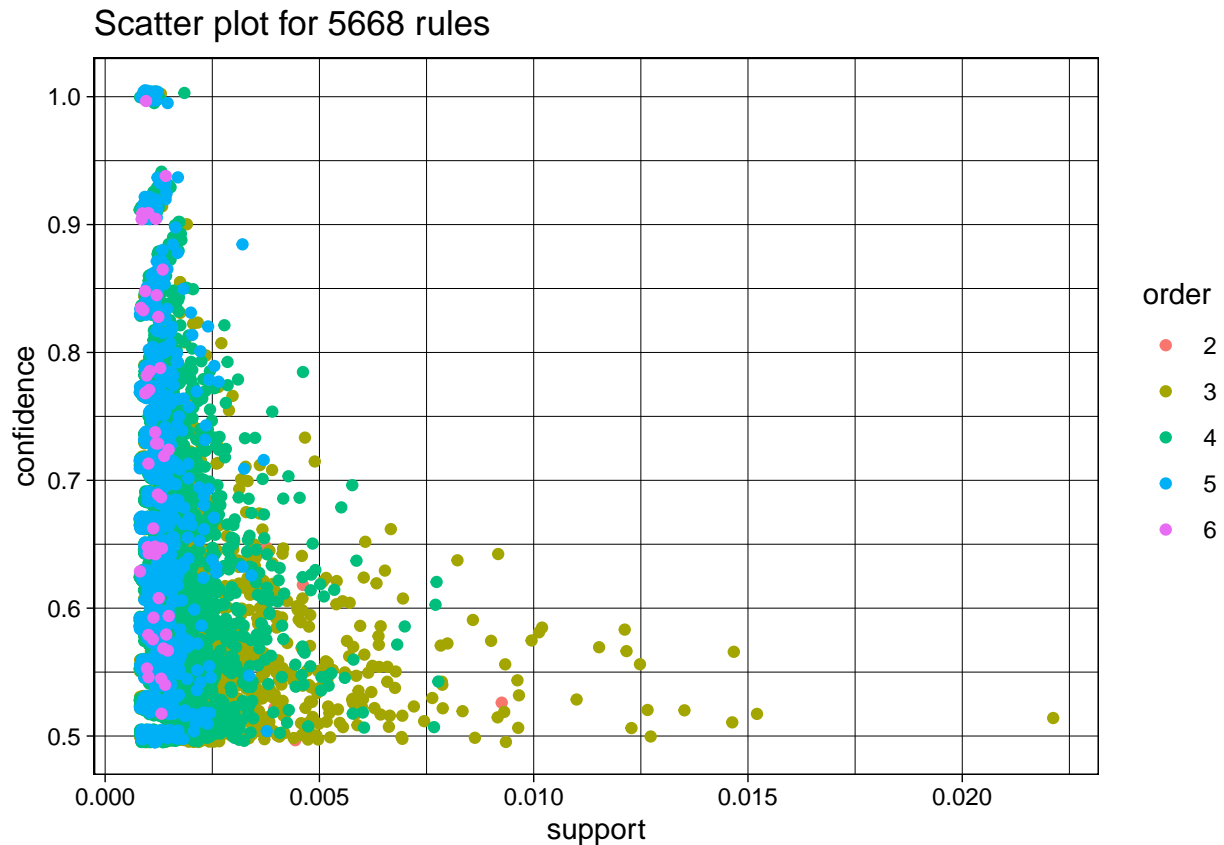
```

##      tropical fruit,
##      yogurt}          => {butter}          0.001118454  0.5238095 0.002135231  9.452599
## [22] {butter,
##      pastry,
##      yogurt}          => {curd}            0.001016777  0.5000000 0.002033554  9.384542
## [23] {butter,
##      hard cheese,
##      whole milk}       => {whipped/sour cream} 0.001423488  0.6666667 0.002135231  9.300236
## [24] {butter,
##      fruit/vegetable juice,
##      other vegetables,
##      tropical fruit}   => {whipped/sour cream} 0.001016777  0.6666667 0.001525165  9.300236
## [25] {citrus fruit,
##      cream cheese,
##      other vegetables,
##      whole milk}       => {domestic eggs}    0.001118454  0.5789474 0.001931876  9.124916
## [26] {cream cheese,
##      curd,
##      whole milk,
##      yogurt}           => {whipped/sour cream} 0.001118454  0.6470588 0.001728521  9.026700
## [27] {hard cheese,
##      other vegetables,
##      tropical fruit}   => {butter}          0.001016777  0.5000000 0.002033554  9.022936
## [28] {napkins,
##      other vegetables,
##      whipped/sour cream,
##      whole milk}       => {butter}          0.001016777  0.5000000 0.002033554  9.022936
## [29] {citrus fruit,
##      cream cheese,
##      whole milk}       => {domestic eggs}    0.001626843  0.5714286 0.002846975  9.006410
## [30] {frozen fish,
##      other vegetables,
##      tropical fruit}   => {pip fruit}        0.001016777  0.6666667 0.001525165  8.812724

```

```
plot(rules, method='two-key plot')
```

```
## To reduce overplotting, jitter is added! Use jitter = 0 to prevent jitter.
```




Two-Key plot shows that baskets with two rules are very rare. It shows a lot of 3, 4, and 5 order rules. 3 order rules tend to have lower values for confidence, but higher levels in support. 4 order rules tend to have a high range of confidence, but low support, indicating they don't appear very often. 5 order rules tend to also have a high range of confidence, but low support.

```
# Load visualization package
library(arulesViz)

# Plot item frequencies
itemFrequencyPlot(transactions, topN = 20, type = "absolute", main = "Top 20 Items by Frequency")
```

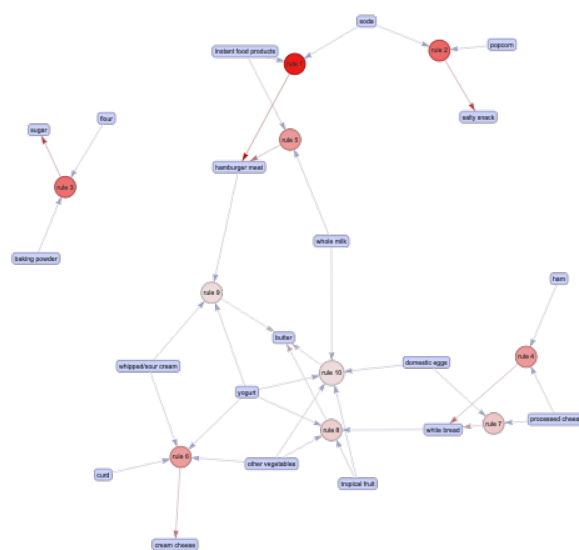
Top 20 Items by Frequency




Plot rules using different methods

```
plot(sorted_rules[1:10], method = "graph", engine = "htmlwidget")
```

Select by id



```
plot(sorted_rules[1:15], method = "grouped")
```



Looking at the item frequency plot, we can see the top items are whole milk, other vegetables, rolls/buns, soda, and yogurt. From here, the count of frequency tends to drop off. These items are likely in a lot of carts, and are likely to have a strong effect on the rules we see. An interesting point on the grouped dot plot shows a point with very high lift and high confidence, which correlates to the relationship between instant food products, soda, and hamburger meat. This may indicate someone who is shopping for cheaper items in the store.

Focusing on High Confidence, High Lift Combination

```
# Find the rules with high confidence and high lift
high_confidence_high_lift <- subset(rules, confidence > 0.8 & lift > 7.5)
arules::inspect(high_confidence_high_lift)
```

##	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{liquor,	=> {bottled beer}	0.001931876	0.9047619	0.002135231	11.235269	19
## [2]	{citrus fruit,	=> {tropical fruit}	0.001118454	0.8461538	0.001321810	8.063879	11
## [3]	{other vegetables,	=> {root vegetables}	0.001321810	0.8666667	0.001525165	7.951182	13
## [4]	{ham,						

```

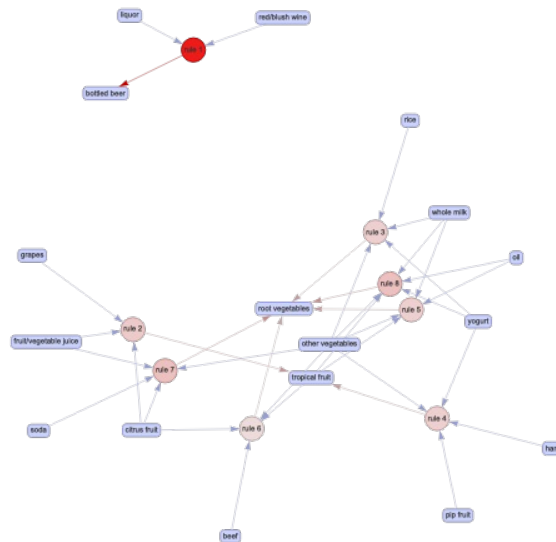
##      other vegetables,
##      pip fruit,
##      yogurt}          => {tropical fruit}  0.001016777  0.8333333  0.001220132  7.941699  10
## [5] {oil,
##      other vegetables,
##      tropical fruit,
##      whole milk}      => {root vegetables} 0.001321810  0.8666667  0.001525165  7.951182  13
## [6] {beef,
##      citrus fruit,
##      other vegetables,
##      tropical fruit}   => {root vegetables} 0.001016777  0.8333333  0.001220132  7.645367  10
## [7] {citrus fruit,
##      fruit/vegetable juice,
##      other vegetables,
##      soda}            => {root vegetables} 0.001016777  0.9090909  0.001118454  8.340400  10
## [8] {oil,
##      other vegetables,
##      tropical fruit,
##      whole milk,
##      yogurt}          => {root vegetables} 0.001016777  0.9090909  0.001118454  8.340400  10

```

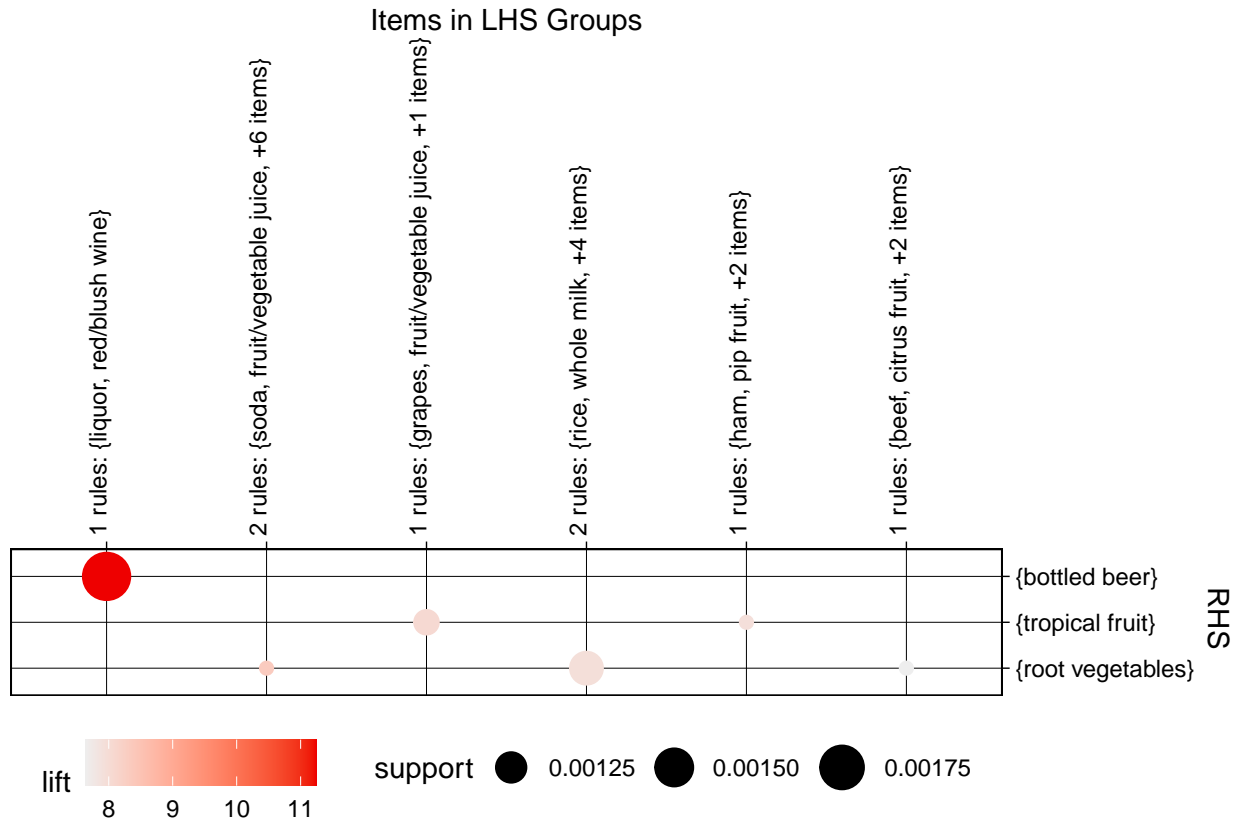
visualizing these rules

```
plot(high_confidence_high_lift, method = "graph", engine = "htmlwidget")
```

Select by id ▾



```
plot(high_confidence_high_lift[1:8], method = "grouped")
```



When we look at the rules with high confidence (>0.8) and high lift (>7.5) only 8 rules show up. In this case, we opted to choose high values of both to see if we could understand any deep patterns between items. When we evaluate the plot above, it seems like the high confidence, high lift, dots tend to correlate to a “general” shop, where customers seem to be buying the basic items (fruit, vegetables, meat, etc).

Interesting Rules to Note:

1. Alcohol Restock : (Highest Confidence and Lift!) - Liquor, Red/blush wine -> Bottled beer - This rule has a high confidence of 0.9 and a lift of 11.4. This means that if a customer buys liquor and red/blush wine, they are 11.4 times more likely to buy bottled beer as well. Likely customers restocking on alcohol or throwing a party!
2. Fruit Salad: Citrus fruit, fruit/vegetable juice, grapes -> tropical fruit - This rule has a high confidence of 0.85 and a lift of 8.5. This means that if a customer buys citrus fruit, fruit/vegetable juice, and grapes, they are 8.5 times more likely to buy tropical fruit as well. Likely customers buying a variety of fruits for a fruit salad or smoothie.
3. Casual Shop: Oil, other vegetables, tropical fruit, whole milk, yogurt -> root vegetables - This rule has a high confidence of 0.91 and a lift of 8.34. This grouping is likely customers shopping casually and buying the essentials for the week.

Question: Why did you choose the thresholds you did?

Thresholds Chosen:

- Support: 0.001 -> Originally tried a support level of 0.01, but this was too high and resulted in only 150 rules being generated. This support level of 0.001 generated 5000 rules when the only limiters were a support of 0.001 and confidence of 0.5.
- Lift: 7.5 -> Originally tried a lift of 1, but rules that seemed to be more “important” had very high lift values, anywhere from 8-22.

- Confidence: 0.5 → Started with a general confidence choice of 0.5, but then used a subset of a confidence level of 0.9 to find the rules with the highest confidence and highest lift.

Question: Do your discovered item sets make sense?

Yes, the item sets that had the highest confidence and lift made sense in terms of the items grouped together. The item sets tend to be items with “interaction” such as if someone wants to make a sandwich, they are very likely to buy bread, cheese, and meat all together. Or if someone is buying a lot of fruits, they are likely to buy a variety of fruits to make a fruit salad or smoothie. If someone is hosting a party, they would tend to buy a variety of alcohol. These item sets make sense in terms of the context of the items grouped together.